

Computer Vision

7. Image Transformation, Matching and Panorama

I-Chen Lin

College of Computer Science
National Yang Ming Chiao Tung University

Outline

- ▶ Image transformation
- ▶ How to represent the environment with a few images?
 - ▶ Key point matching.
 - ▶ Automatic image stitching.

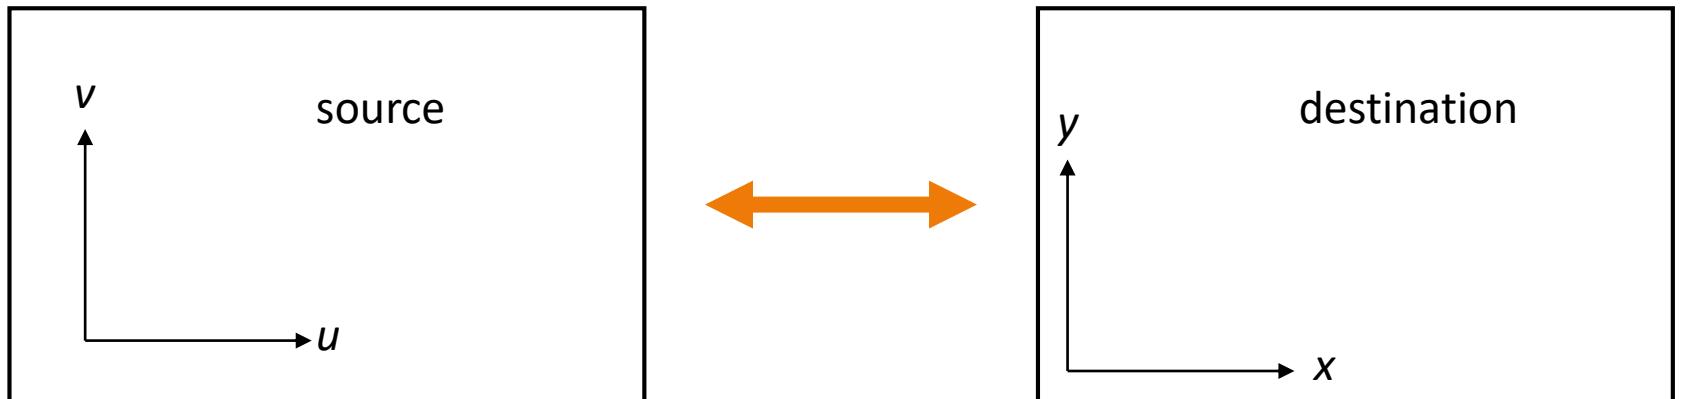
Several figures are from the following reference lists:

- Image-based Modeling and Rendering, SIGGRAPH'99 course notes.
- L. McMillan, G. Bishop, "Plenoptic Modeling: An Image-Based Rendering System", Proc. SIGGRAPH'95, pp. 39-46.
- S.E. Chen, "QuickTime VR – An Image-Based Approach to Virtual Environment Navigation", Proc. SIGGRAPH'95, pp. 29-38.
- M. Brown and D.G. Lowe, "Recognising panoramas," Proc., ICCV 2003, pp. 1218-25.
- M. Brown and D.G. Lowe, "Automatic Panoramic Image Stitching using Invariant Features," Intl. J. Computer Vision, vol.74, no.1, pp.59-73.
- R. Szeliski, "Image Alignment and Stitching: A Tutorial," Foundations and Trends in Computer Graphics and Vision, vol.2, no.1., pp.1-104.

Image warping and transformation

- ▶ Rearranging pixels of a picture.
 - ▶ It's useful for both image processing and for computer graphics (namely, for texture mapping).
- ▶ Finding corresponding points in the source and destination images.
- ▶ This function is called the “mapping” or “transformation”.

Image warping (cont.)



Forward mapping : $(x, y) = f(u, v)$

Inverse mapping : $(u, v) = f'(x, y)$

Mapping types

- ▶ How to create the mapping systematically?
- ▶ Simple mappings
 - ▶ affine mapping
 - ▶ projective mapping
 - ▶ bilinear mapping
 - ▶

Affine mappings

- ▶ $u = ax + by + c$
- ▶ $v = dx + ey + f$

$$\begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \quad \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

- ▶ A combination of 2-D scale, rotation, and translation transformations.
- ▶ Allows a square to be distorted into any parallelogram.
- ▶ 6 degrees of freedom.
- ▶ Inverse is of same form (is also affine).

Projective mappings

- ▶ $u = (ax+by+c)/(gx+hy+i)$
- ▶ $v = (dx+ey+f)/(gx+hy+i)$

$$\begin{bmatrix} uq \\ vq \\ q \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

$$u = uq/q, v = vq/q$$

- ▶ Linear numerator and denominator.
- ▶ If $g=h=0$ then you get affine as a special case.
- ▶ Allows a square to be distorted into any quadrilateral.
- ▶ 8 degrees of freedom.
- ▶ Inverse is of same form (is also projective)

Projective mappings

- ▶ Every point correspondence can have two equations:

- ▶ $x(XA_{31} + YA_{32} + A_{33}) = (XA_{11} + YA_{12} + A_{13})$
- ▶ $y(XA_{31} + YA_{32} + A_{33}) = (XA_{21} + YA_{22} + A_{23})$

$$\begin{bmatrix} wx \\ wy \\ w \end{bmatrix} = \begin{bmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{bmatrix} * \begin{bmatrix} X \\ Y \\ 1 \end{bmatrix}$$

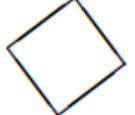
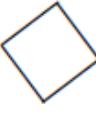
- ▶ After rearrangement, a pair of quadrilaterals:

$$\begin{array}{ccccccccc} X_1 & Y_1 & 1 & 0 & 0 & 0 & -x_1X_1 & -x_1Y_1 & -x_1 \\ 0 & 0 & 0 & X_1 & Y_1 & 1 & -y_1X_1 & -y_1Y_1 & -y_1 \\ X_2 & Y_2 & 1 & 0 & 0 & 0 & -x_2X_2 & -x_2Y_2 & -x_2 \\ 0 & 0 & 0 & X_2 & Y_2 & 1 & -y_2X_2 & -y_2Y_2 & -y_2 \\ X_3 & Y_3 & 1 & 0 & 0 & 0 & -x_3X_3 & -x_3Y_3 & -x_3 \\ 0 & 0 & 0 & X_3 & Y_3 & 1 & -y_3X_3 & -y_3Y_3 & -y_3 \\ X_4 & Y_4 & 1 & 0 & 0 & 0 & -x_4X_4 & -x_4Y_4 & -x_4 \\ 0 & 0 & 0 & X_4 & Y_4 & 1 & -y_4X_4 & -y_4Y_4 & -y_4 \end{array} * \begin{bmatrix} A_{11} \\ A_{12} \\ A_{13} \\ A_{21} \\ A_{22} \\ A_{23} \\ A_{31} \\ A_{32} \\ A_{33} \end{bmatrix} = 0$$

$\min E = |Ux|^2$, subject to $|x|=1$

The x that minimize E is eigenvector e_1 of $U^T U$

Properties of image transformation

Name	Matrix	Number of d.o.f.	Preserves	Icon
Translation	$[\mathbf{I} \mathbf{t}]_{2 \times 3}$	2	Orientation + ...	
Rigid (Euclidean)	$[\mathbf{R} \mathbf{t}]_{2 \times 3}$	3	Lengths + ...	
Similarity	$[s\mathbf{R} \mathbf{t}]_{2 \times 3}$	4	Angles + ...	
Affine	$[\mathbf{A}]_{2 \times 3}$	6	Parallelism + ...	
Projective	$[\tilde{\mathbf{H}}]_{3 \times 3}$	8	Straight lines	

Performing an image warp

- ▶ Source scanning:
*for v = vmin to vmax
 for u = umin to umax
 $x = x(u,v)$
 $y = y(u,v)$
 copy pixel at source[u,v] to dest[x,y]*
- ▶ Destination scanning:
*for y = ymin to ymax
 for x = xmin to xmax
 $u = u(x,y)$
 $v = v(x,y)$
 copy pixel at source[u,v] to dest[x,y]*

Is there any problem?

Acquiring cylindrical projections

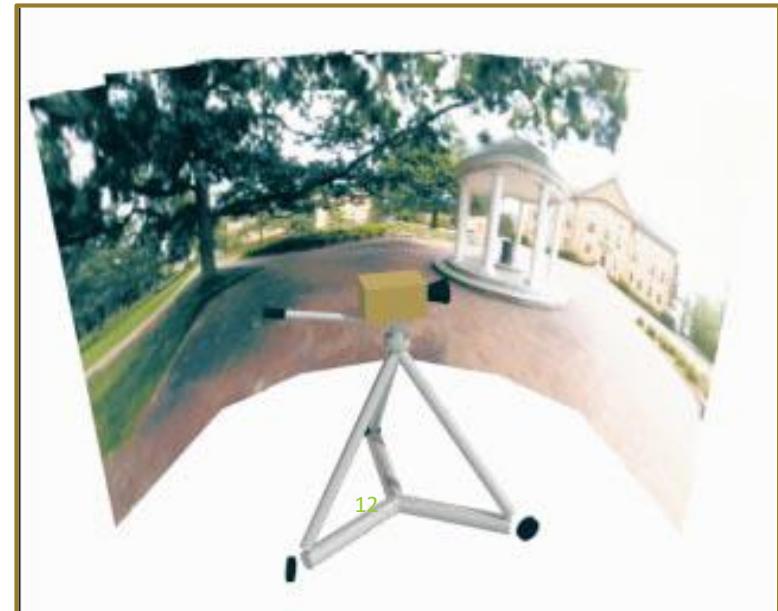


11

S.E. Chen, "QuickTime VR", Proc. SIGGRAPH'95

Acquiring cylindrical projections

- ▶ Using a regular camera with a panning tripod.
- ▶ Camera model and rotations can be inferred from planar reference images. (without camera calibration)
- ▶ Combine planar images and then compute the projection onto a cylinder.



Panoramic images



Perspective projection on view plane



13



Automatic Panoramic Image Stitching using Invariant Features

Proposed by M. Brown and D.G. Lowe,
In Proc., ICCV 2003 and IJCV 2007

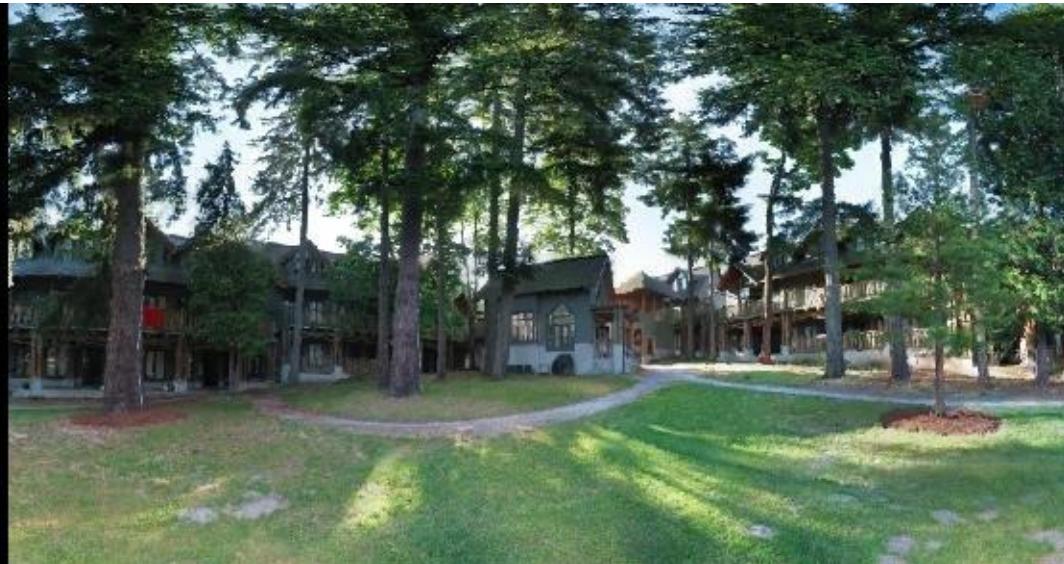
Introduction

- ▶ Are you getting the whole picture?
 - ▶ Compact Camera FOV = $50 \times 35^\circ$



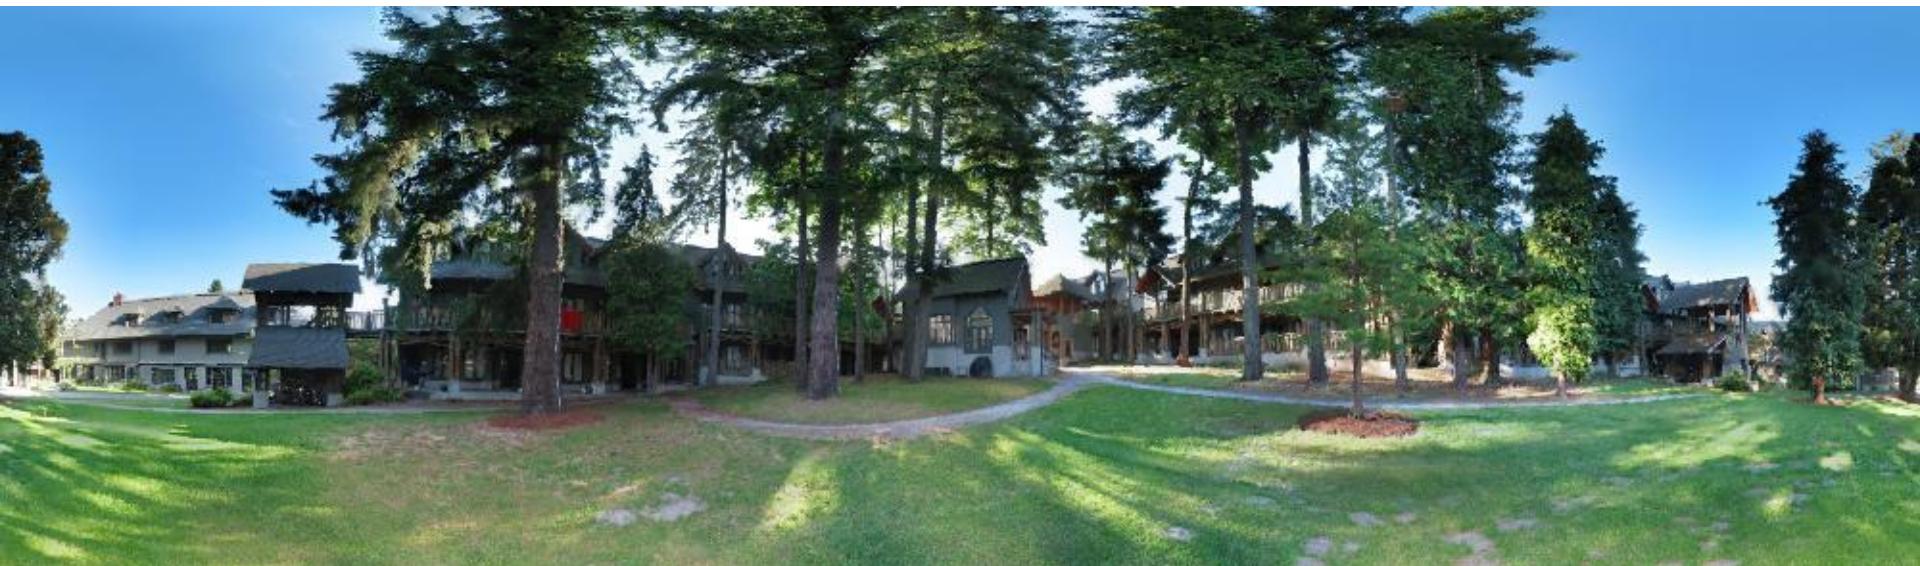
Introduction

- ▶ Are you getting the whole picture?
 - ▶ Compact Camera FOV = $50 \times 35^\circ$
 - ▶ Human FOV = $200 \times 135^\circ$



Introduction

- ▶ Are you getting the whole picture?
 - ▶ Compact Camera FOV = $50 \times 35^\circ$
 - ▶ Human FOV = $200 \times 135^\circ$
 - ▶ Panoramic Mosaic = $360 \times 180^\circ$



Why “Recognising Panoramas”?

Why “Recognising Panoramas”?

- ▶ 1D Rotations (θ)
 - ▶ Ordering \Rightarrow matching images

Why “Recognising Panoramas”?

- ▶ 1D Rotations (θ)
 - ▶ Ordering \Rightarrow matching images



Why “Recognising Panoramas”?

- ▶ 1D Rotations (θ)
 - ▶ Ordering \Rightarrow matching images



Why “Recognising Panoramas”?

- ▶ 1D Rotations (θ)
 - ▶ Ordering \Rightarrow matching images



- 2D Rotations (θ, ϕ)
 - Ordering $\not\Rightarrow$ matching images

Why “Recognising Panoramas”?

- ▶ 1D Rotations (θ)
 - ▶ Ordering \Rightarrow matching images



- 2D Rotations (θ, ϕ)
 - Ordering $\not\Rightarrow$ matching images



Why “Recognising Panoramas”?

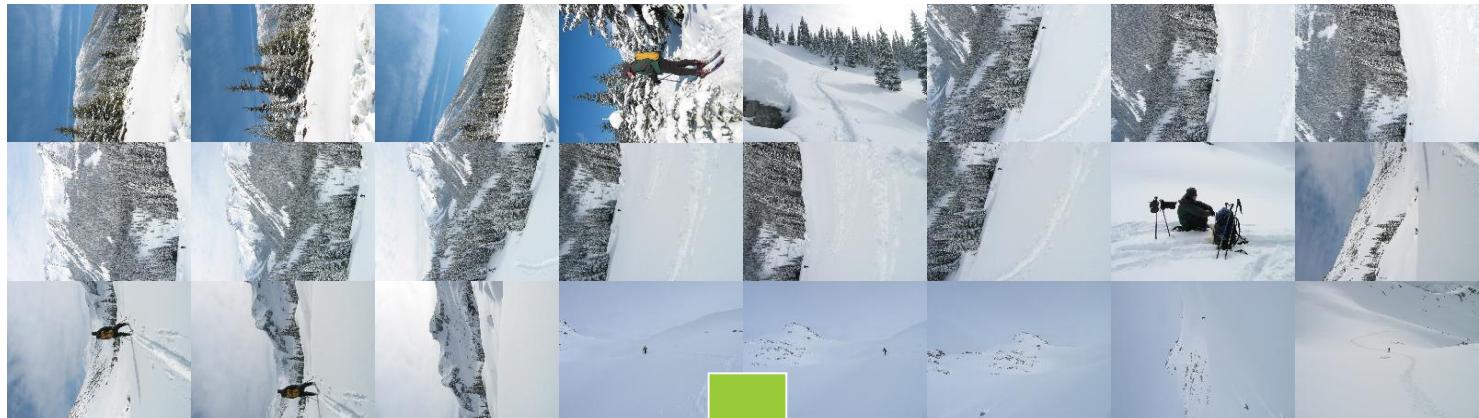
- ▶ 1D Rotations (θ)
 - ▶ Ordering \Rightarrow matching images



- 2D Rotations (θ, ϕ)
 - Ordering $\not\Rightarrow$ matching images



Why “Recognising Panoramas”?



Overview

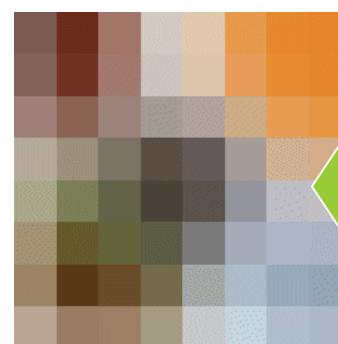
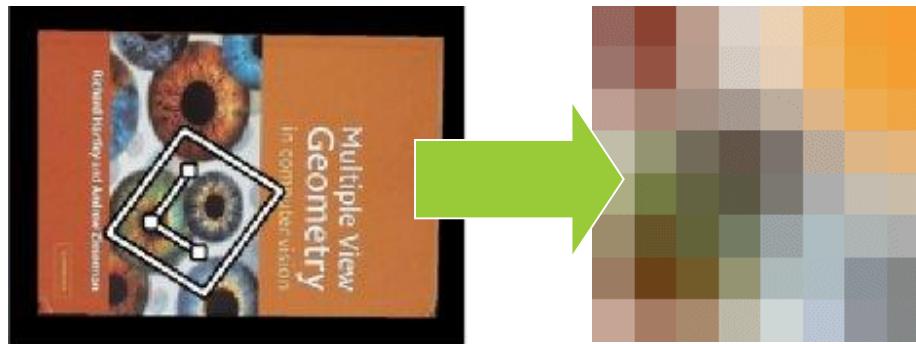
- ▶ Feature Matching
- ▶ Image Matching
- ▶ Bundle Adjustment
- ▶ Multi-band Blending
- ▶ Results
- ▶ Conclusions

Overview

- ▶ Feature Matching
 - ▶ SIFT Features
 - ▶ Nearest Neighbor Matching
- ▶ Image Matching
- ▶ Bundle Adjustment
- ▶ Multi-band Blending
- ▶ Results
- ▶ Conclusions

Invariant features

- ▶ Schmid & Mohr 1997, Lowe 1999, Baumberg 2000, Tuytelaars & Van Gool 2000, Mikolajczyk & Schmid 2001, Brown & Lowe 2002, Matas et. al. 2002, Schaffalitzky & Zisserman 2002

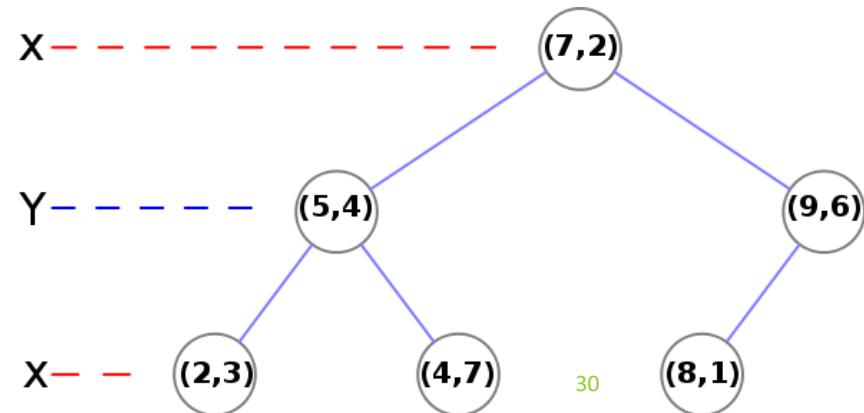
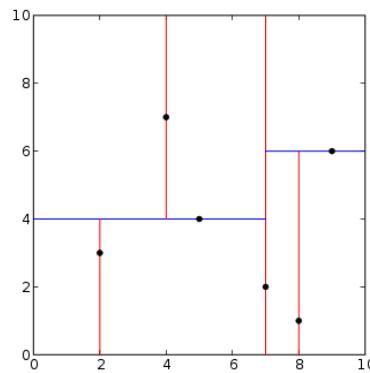


SIFT features

- ▶ Invariant Features
 - ▶ Establish invariant frame
 - ▶ Maxima/minima of scale-space DOG $\Rightarrow x, y, s$
 - ▶ Maximum of distribution of local gradients $\Rightarrow \theta$
 - ▶ Form descriptor vector
 - ▶ Histogram of smoothed local gradients
 - ▶ 128 dimensions
- ▶ SIFT features are...
 - ▶ Geometrically invariant to similarity transforms,
 - ▶ some robustness to affine change
 - ▶ Photometrically invariant to affine changes in intensity

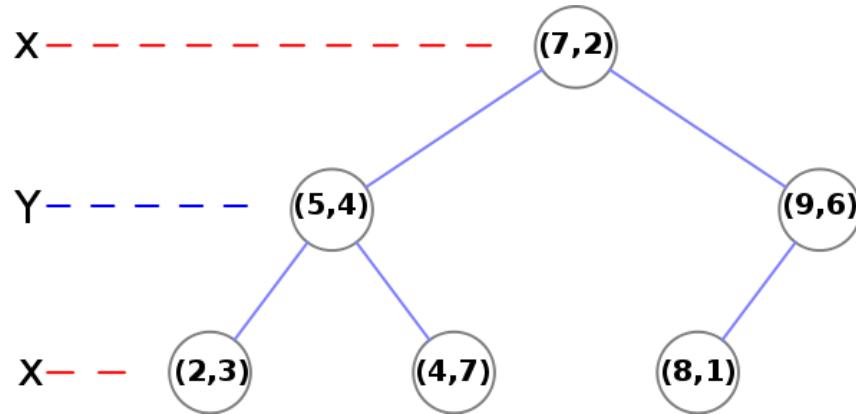
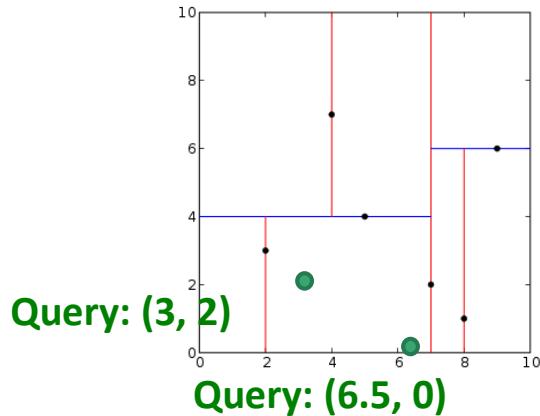
Nearest neighbor matching

- ▶ Find k-NN for each feature
 - ▶ $k \approx$ number of overlapping images (in this work, $k = 4$)
- ▶ Use k-d tree or other searching tech. for speed-up
 - ▶ k-d tree recursively bi-partitions data at mean in the dimension of maximum variance.
 - ▶ Tree building complexity is less than $O(kn\log n)$.



Figures from: <http://en.wikipedia.org/wiki/Kd-tree>

k-d tree



- In a recursive style
 - Top-down search to a leaf node.
 - Backward check whether there are “possible” NN through the distance to the hyperplane of the parents’ nodes.
- For a large K(dimension), approximate nearest neighbors are popularly used.

Overview

- ▶ Feature Matching
- ▶ Image Matching
 - ▶ RANSAC for Homography
- ▶ Bundle Adjustment
- ▶ Multi-band Blending
- ▶ Results
- ▶ Conclusions

Matching with features

- ▶ For each keypoint, we have multiple candidates.
- ▶ How to correctly recognize the corresponding one?



RANSAC for point matching

- ▶ If there're multiple candidates, "*RAN*dom *SAmple Consensus*" is an approach to solve the ambiguity.
- ▶ The basic idea is to sample a set of point correspondences.
- ▶ Use the sample to estimate a motion model (mapping between 2 views)
 - ▶ E.g. translation, fundamental matrix, etc.
- ▶ See whether the remaining points can provide support for this solution.
- ▶ Copes with a large proportion of outliers.

RANSAC example

- ▶ E.g. fitting a straight line

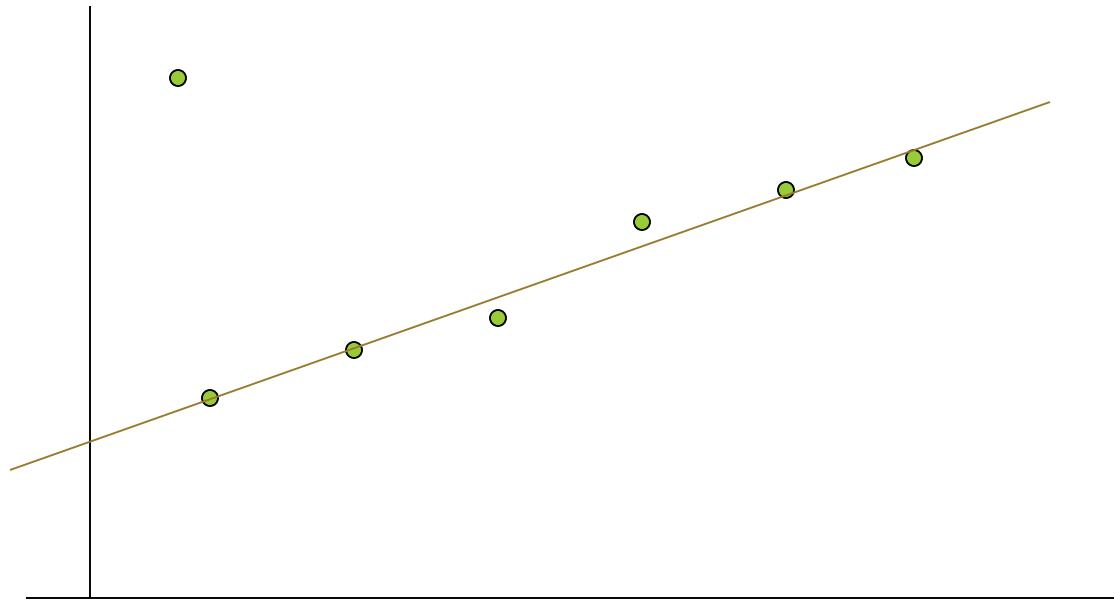
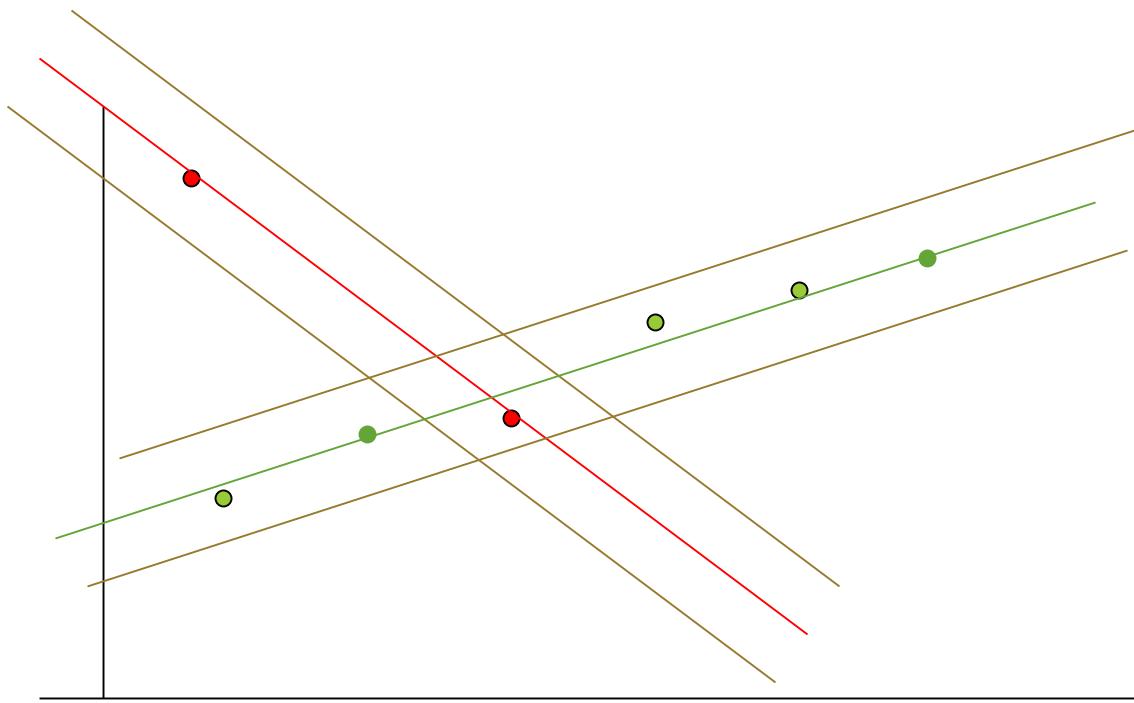


Figure from Prof. J. Rehg, Computer Vision, Georgia Inst. of Tech.

Main idea

- ▶ Select 2 points at random
- ▶ Fit a line
- ▶ “Support” = number of inliers
- ▶ Line with most inliers wins

Why will this work ?



The best line has the most support
More support -> Better fit

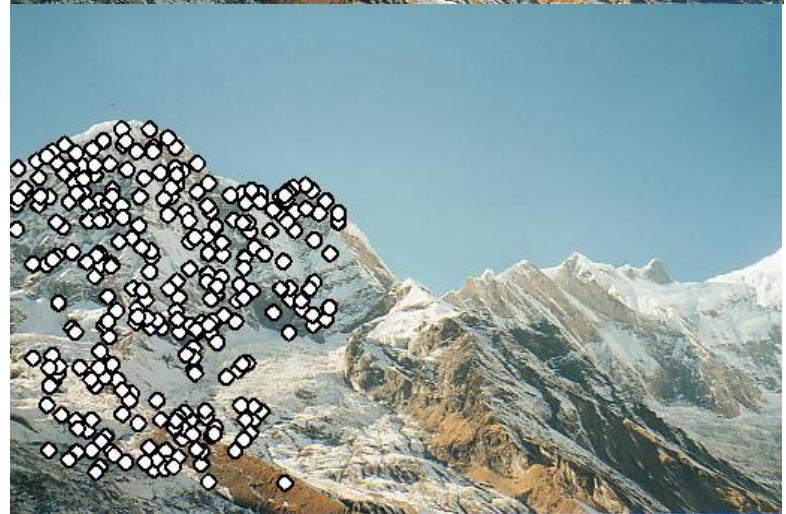
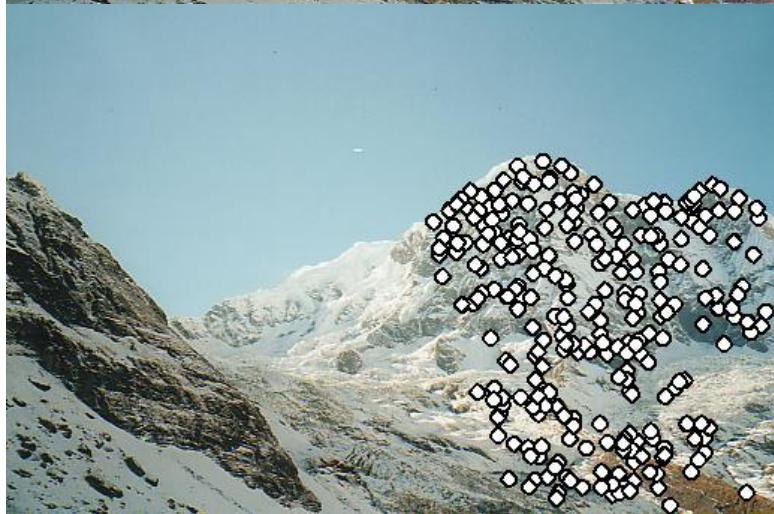
RANSAC

- ▶ Objective:
 - ▶ Robust fit of a model to data S
- ▶ Algorithm
 - ▶ Randomly select s points
 - ▶ Instantiate a model
 - ▶ Get consensus set S_i
 - ▶ If $|S_i| > T$, terminate and return model
 - ▶ Repeat for N trials, return model with max $|S_i|$

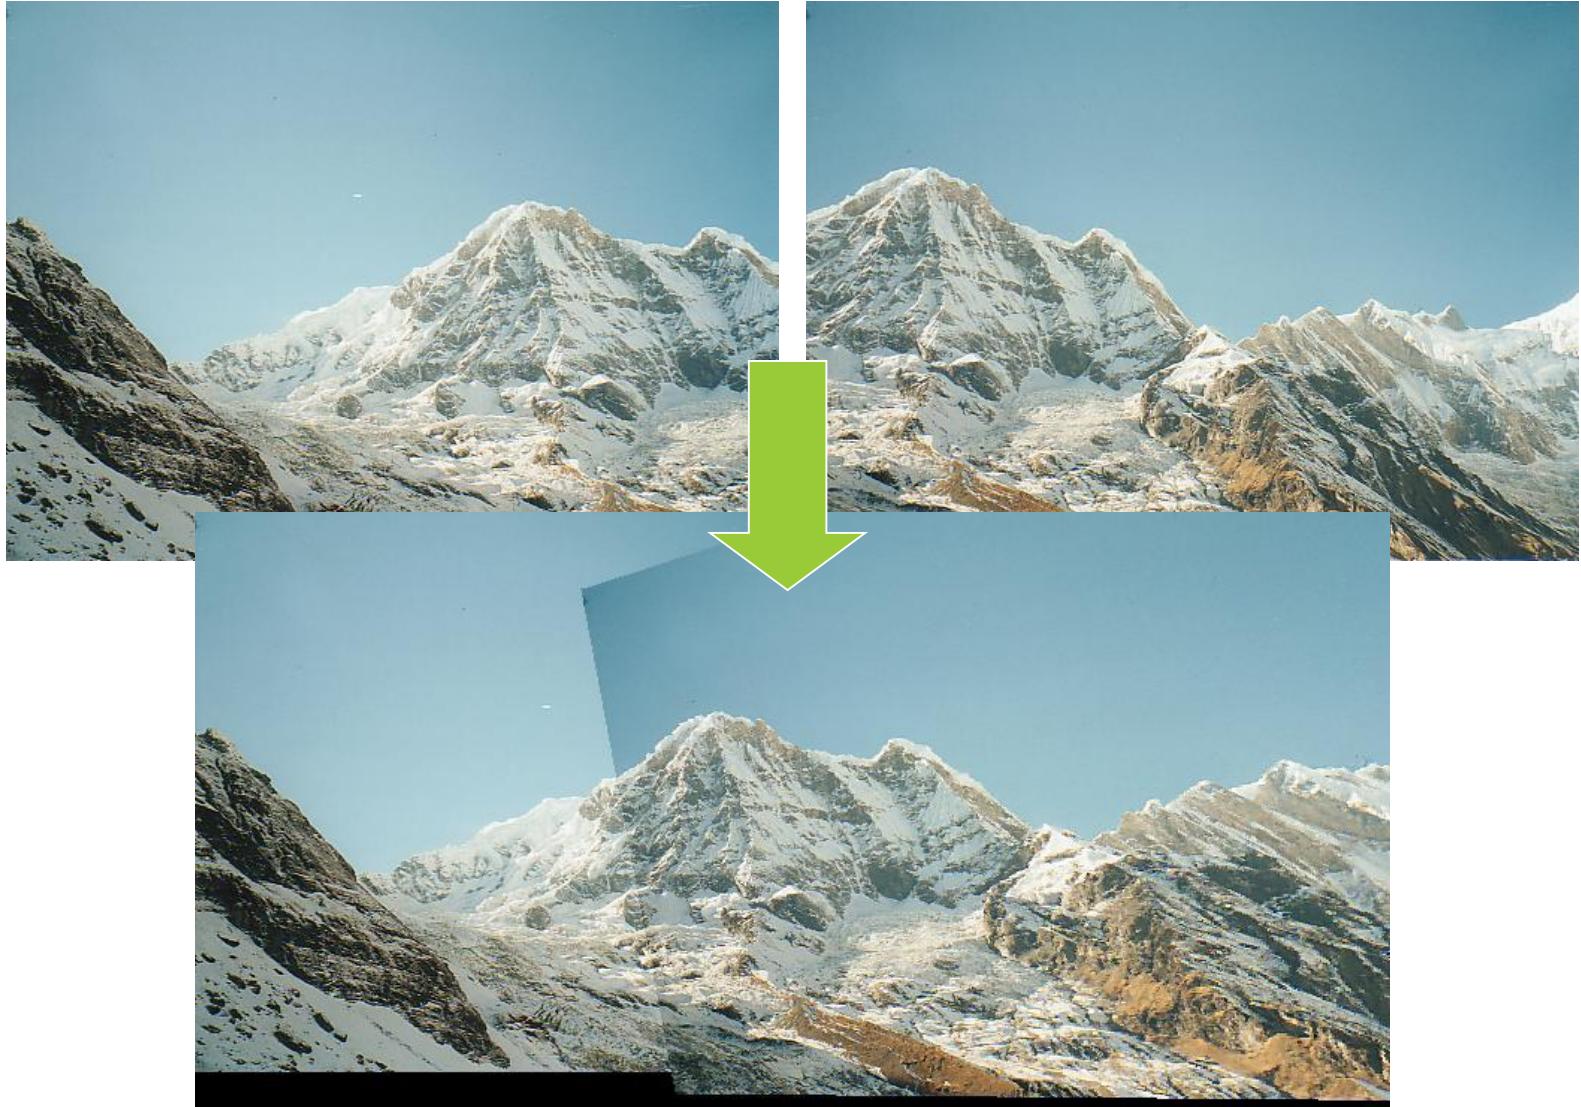
RANSAC for homography



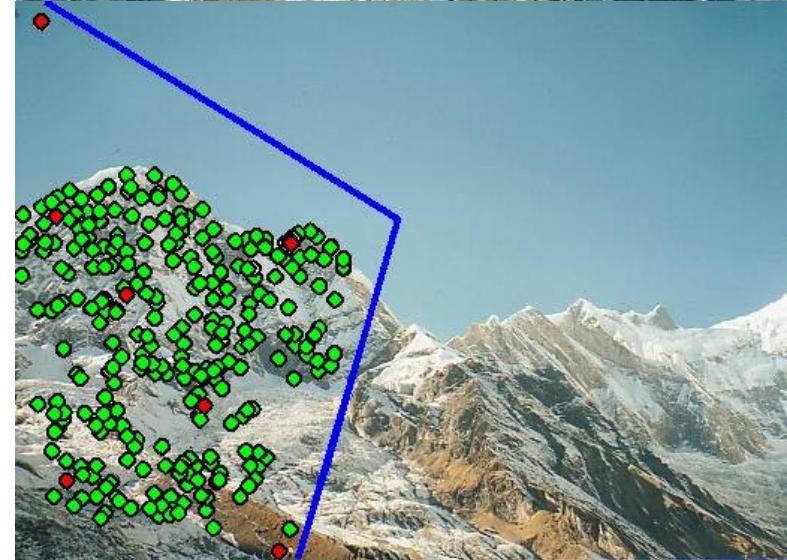
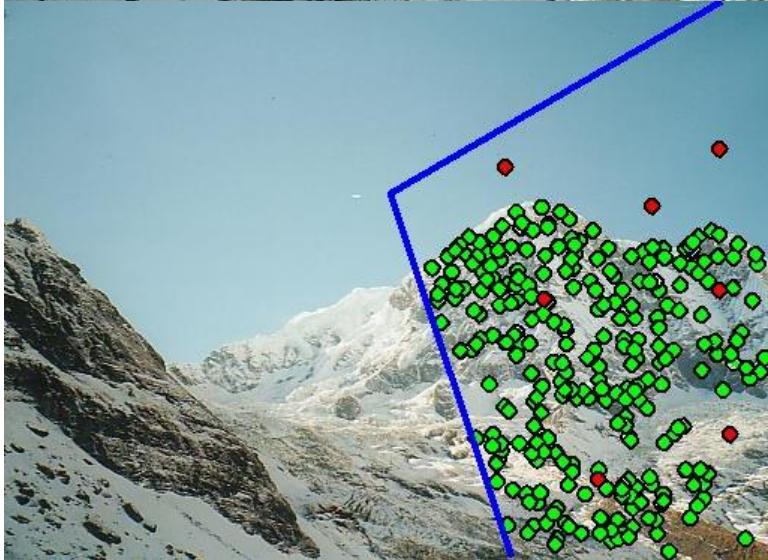
RANSAC for homography



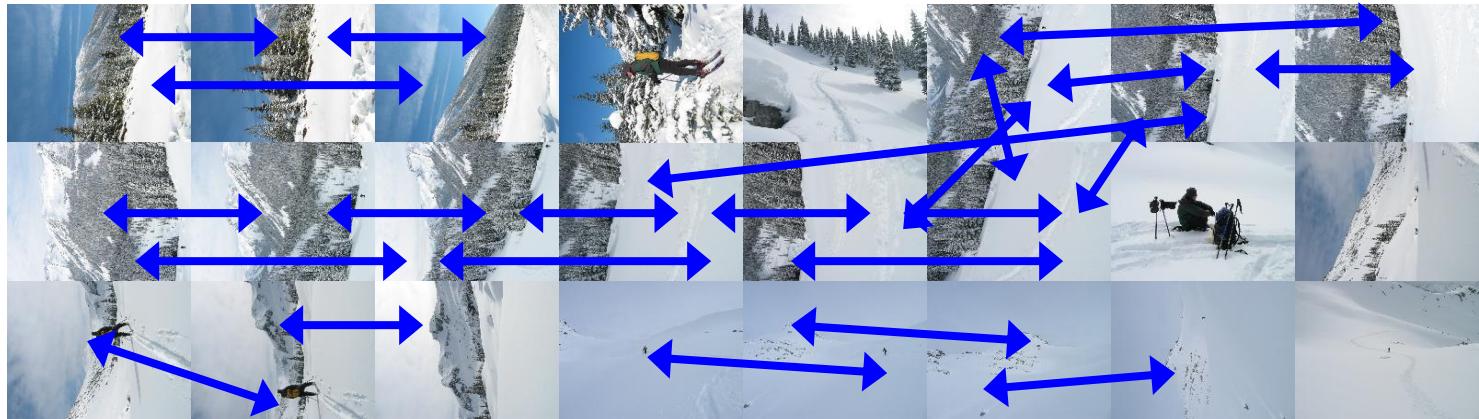
RANSAC for homography



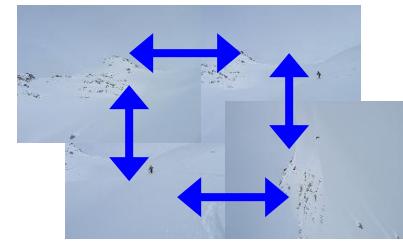
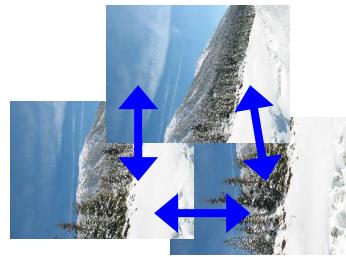
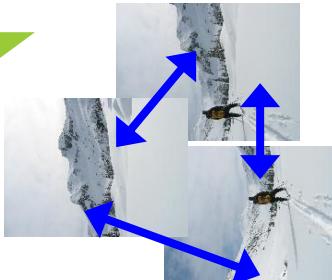
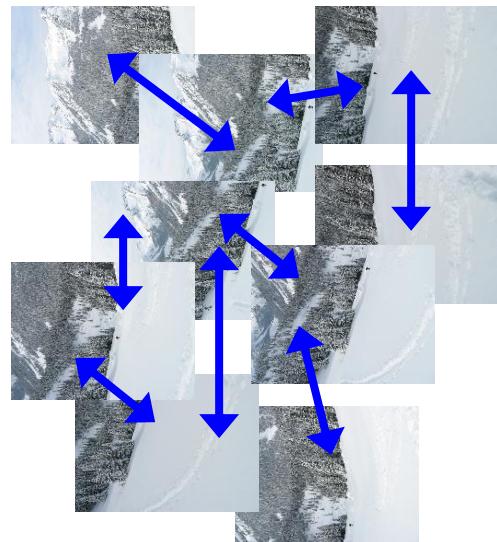
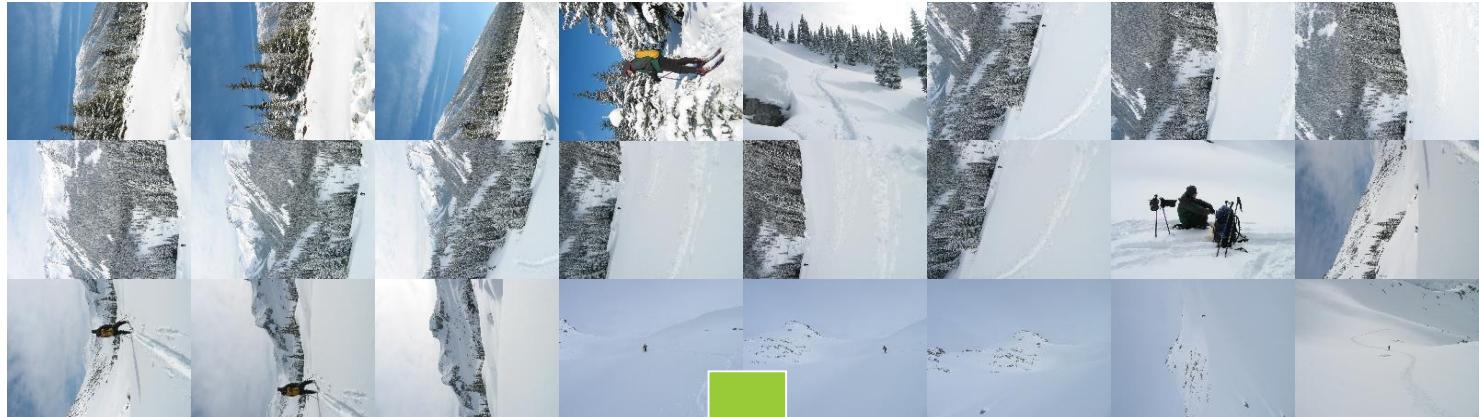
Probabilistic model for verification



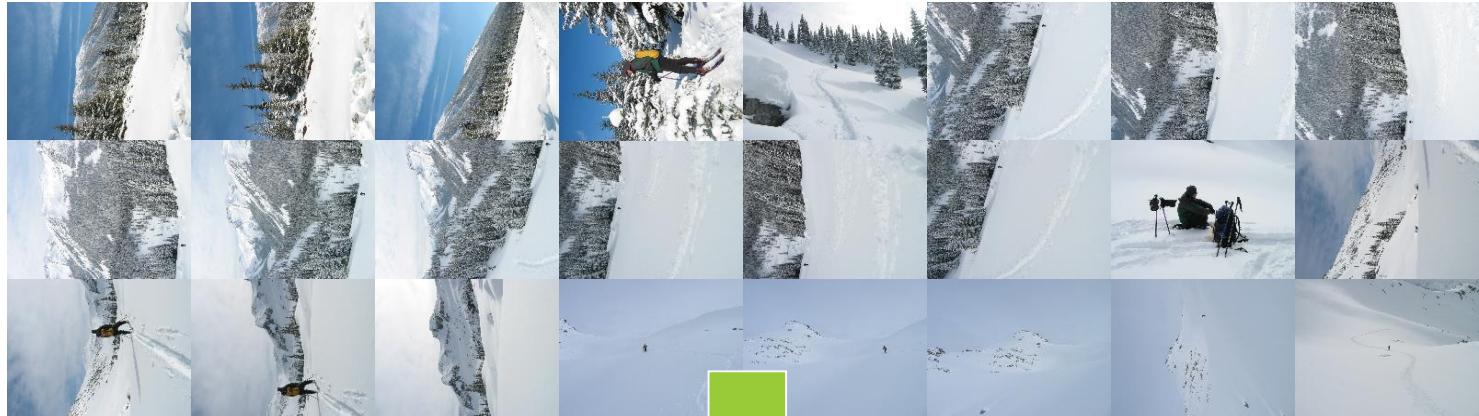
Finding the panoramas (match 6 images at most)



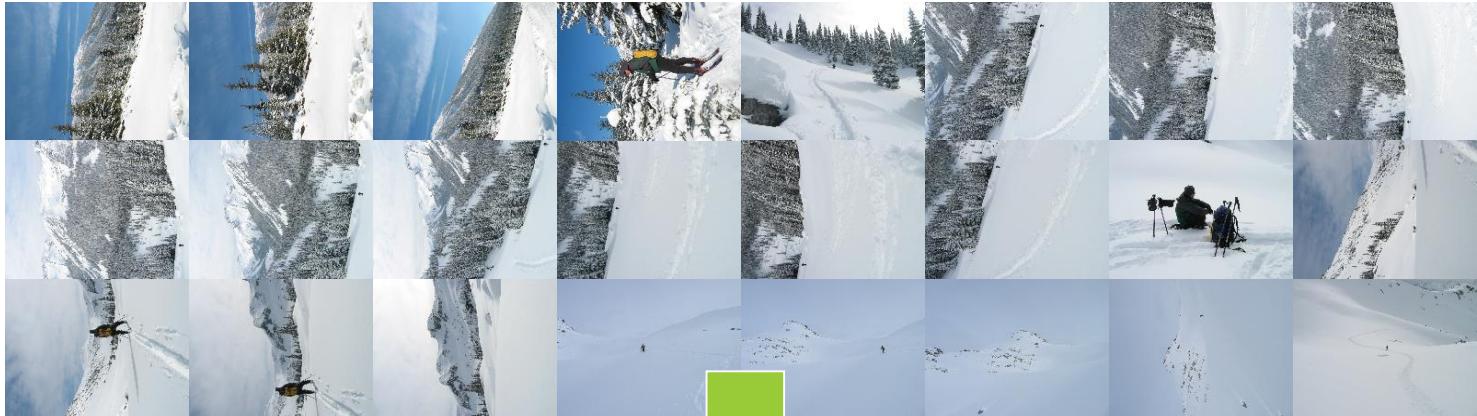
Finding the panoramas



Finding the panoramas



Finding the panoramas



Overview

- ▶ Feature Matching
- ▶ Image Matching
- ▶ Bundle Adjustment
 - ▶ Error function
- ▶ Multi-band Blending
- ▶ Results
- ▶ Conclusions

Error function

f: Require knowledge about camera and R, T

Let's skip the detail of this part !

- ▶ Sum of squared projection errors

$$e = \sum_{i=1}^n \sum_{j \in \mathcal{I}(i)} \sum_{k \in \mathcal{F}(i,j)} f(\mathbf{r}_{ij}^k)^2$$

- ▶ n = #images
- ▶ $\mathcal{I}(i)$ = set of image matches to image i
- ▶ $\mathcal{F}(i, j)$ = set of feature matches between images i,j
- ▶ \mathbf{r}_{ij}^k = residual of kth feature match between images i,j

- ▶ Robust error function

$$f(\mathbf{x}) = \begin{cases} |\mathbf{x}|, & \text{if } |\mathbf{x}| < x_{max} \\ x_{max}, & \text{if } |\mathbf{x}| \geq x_{max} \end{cases}$$

Bundle adjustment

- ▶ New images initialised with rotation, focal length of best matching image.



Bundle adjustment

- ▶ New images initialised with rotation, focal length of best matching image.



Gain compensation

N_{ij} : the number of pixels in image i that overlap in image j



- ▶ Find the optimize gains of g_i according to means of overlapping regions between image pair i and j

$$e = \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^n N_{ij} ((g_i \bar{I}_{ij} - g_j \bar{I}_{ji})^2 / \sigma_N^2 + (1 - g_i)^2 / \sigma_g^2)$$



Overview

- ▶ Feature Matching
- ▶ Image Matching
- ▶ Bundle Adjustment
- ▶ Multi-band Blending
- ▶ Results
- ▶ Conclusions

Multi-band blending

- ▶ Burt & Adelson 1983
 - ▶ Blend frequency bands over range $\propto \lambda$



Linear blending

- ▶ View-dependent weight:

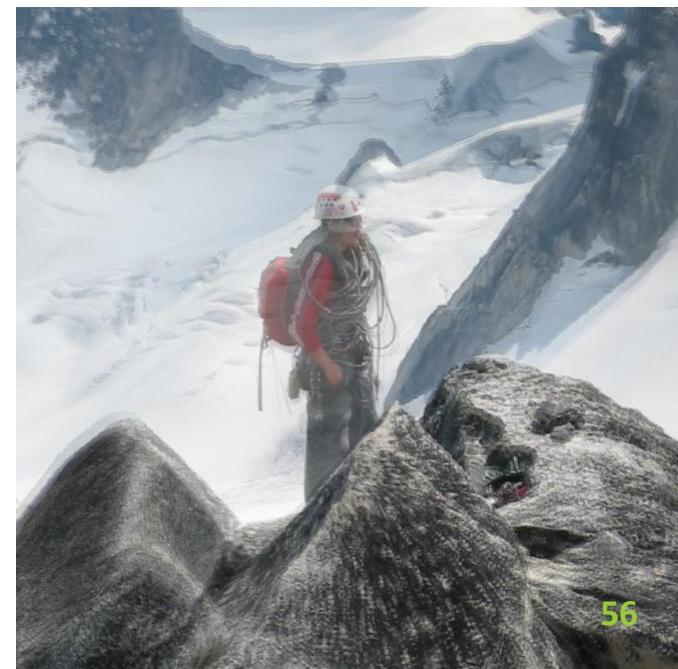
$$W(x, y) = w(x)w(y)$$

where $w(x)$ varies linearly from 1 at the centre of the image to 0 at the edge

- ▶ Linear blending by the weights:

$$I^{linear}(\theta, \phi) = \frac{\sum_{i=1}^n I^i(\theta, \phi)W^i(\theta, \phi)}{\sum_{i=1}^n W^i(\theta, \phi)}$$

- ▶ Cause blurring of high frequency detail if there are small registration errors



Multi-band blending

- ▶ Mainly use the most frontal pixels.

$$W_{max}^i(\theta, \phi) = \begin{cases} 1 & \text{if } W^i(\theta, \phi) = \arg \max_j W^j(\theta, \phi) \\ 0 & \text{otherwise} \end{cases}$$

- ▶ Propagate the weights to neighbors by Gaussian filter.

$$W_\sigma^i(\theta, \phi) = W_{max}^i(\theta, \phi) * g_\sigma(\theta, \phi)$$

- ▶ Divide the image into multiple bands.

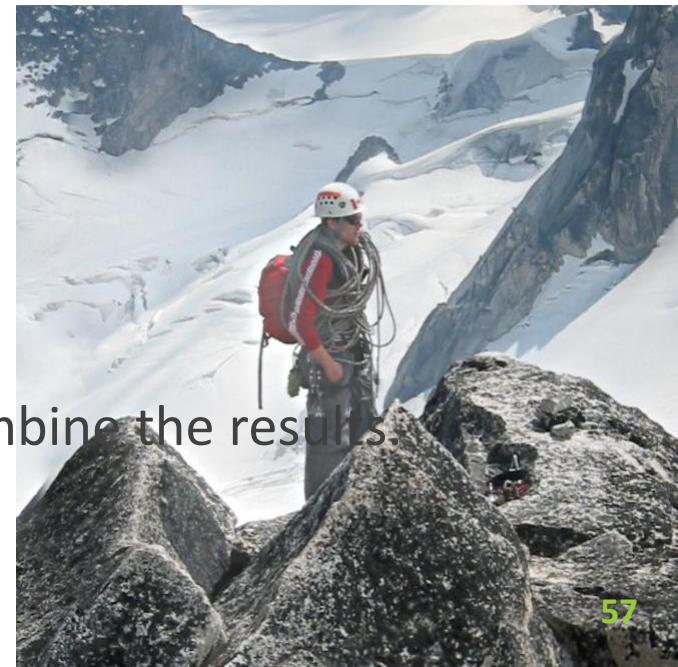
$$B_{(k+1)\sigma}^i = I_{k\sigma}^i - I_{(k+1)\sigma}^i$$

$$I_{(k+1)\sigma}^i = I_{k\sigma}^i * g_{\sigma'}$$

$$W_{(k+1)\sigma}^i = W_{k\sigma}^i * g_{\sigma'}$$

- ▶ Apply linear blending on each band, and combine the results.

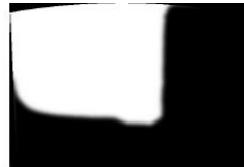
$$I_{k\sigma}^{multi}(\theta, \phi) = \frac{\sum_{i=1}^n B_{k\sigma}^i(\theta, \phi) W_{k\sigma}^i(\theta, \phi)}{\sum_{i=1}^n W_{k\sigma}^i(\theta, \phi)}$$



Multi-band blending



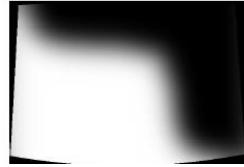
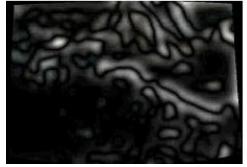
Multi-band blending



Band 1 scale 0 to σ



Band 2 scale σ to 2σ



Band 3 lower than 2σ



Results



Conclusions

- ▶ Fully automatic panoramas
 - ▶ A recognition problem...
- ▶ Invariant feature based method
 - ▶ SIFT features, RANSAC, Bundle Adjustment, Multi-band Blending
 - ▶ $O(n \log n)$
- ▶ Future work of this article
 - ▶ Advanced camera modelling
 - ▶ radial distortion, camera motion, scene motion, vignetting, exposure, high dynamic range, flash ...
 - ▶ Full 3D case – recognising 3D objects/scenes in unordered datasets